

10707

Deep Learning: Spring 2021

Andrej Risteski

Machine Learning Department

Lecture 2:

Representational power of
neural networks

Neural network basics: the artificial neuron

Neuron **pre-activation**:

$$a(\mathbf{x}) = b + \sum_i w_i x_i = b + \mathbf{w}^T \mathbf{x}$$

Neuron **post-activation**:

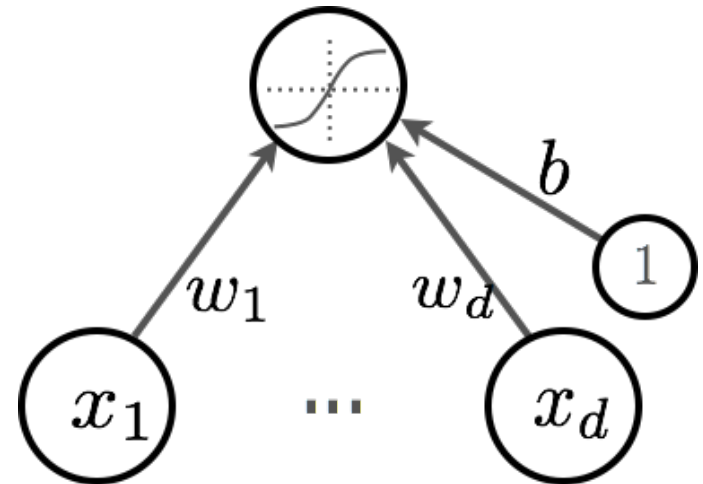
$$h(\mathbf{x}) = \sigma(b + \mathbf{w}^T \mathbf{x})$$

Where:

\mathbf{w} are the **weights** (parameters)

b is the **bias** term

$\sigma(\cdot)$ is called the **activation function**

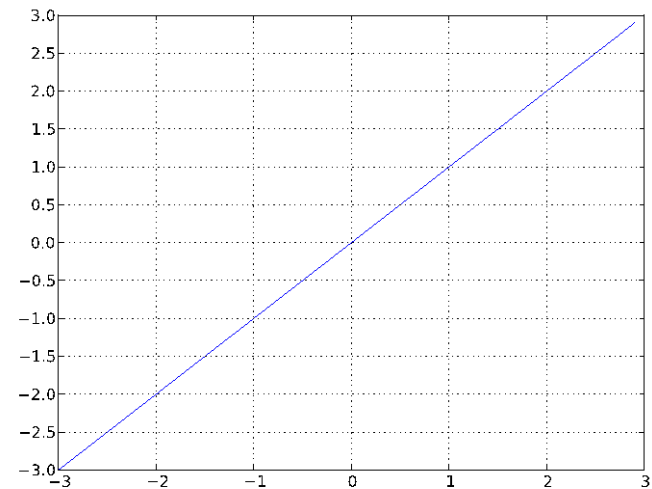


Popular activations

Linear activation function:

$$\sigma(a) = a$$

- ⌘ No nonlinear transformation
- ⌘ No output squashing
- ⌘ Poor representational power (linear composed w/ linear = linear)

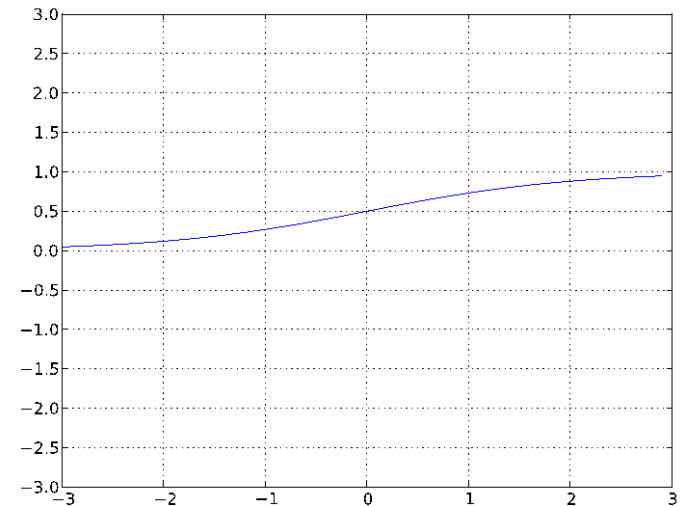


Popular activations

Sigmoid activation function:

$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$

- ⌘ Squashes the neuron's output between 0 and 1: can be interpreted as $P(\text{output} = 1|a)$ (i.e. **logistic classifier**)
- ⌘ Always positive
- ⌘ Bounded
- ⌘ Strictly Increasing




Aside: classification power of a single neuron

Sigmoid activation function: $\sigma(a) = \frac{1}{1 + \exp(-a)}$

- ⊗ Squashes the neuron's output between 0 and 1: can be interpreted as $P(\text{output} = 1|a)$ (i.e. **logistic classifier**)



If activation is greater than 0.5, predict 1. Otherwise predict 0



Can perfectly classify linearly separable datasets, e.g. OR, AND,


Aside: classification power of a single neuron

Sigmoid activation function: $\sigma(a) = \frac{1}{1 + \exp(-a)}$

- ⌘ Squashes the neuron's output between 0 and 1: can be interpreted as $P(\text{output} = 1|a)$ (i.e. **logistic classifier**)



If activation is greater than 0.5, predict 1. Otherwise predict 0



Cannot perfectly classify linearly non-separable datasets, e.g. XOR.

('69, Minsky and Papert, *Perceptrons*)


Popular activations

Hyperbolic tangent (“tanh”) activation function:

$$\sigma(a) = \tanh(a)$$

$$= \frac{\exp(a) - \exp(-a)}{\exp(a) + \exp(-a)} = \frac{\exp(2a) - 1}{\exp(2a) + 1}$$

- ⌘ Squashes neuron's output between -1 and 1
- ⌘ Can be positive or negative
- ⌘ Bounded
- ⌘ Strictly increasing




Popular activations

Rectified linear (“ReLU”) activation function:

$$\sigma(a) = \max(a, 0)$$

- ⌘ Bounded below by 0 (always non-negative)
- ⌘ Tends to produce units with sparse activities
- ⌘ Not upper bounded
- ⌘ Strictly increasing



Single Hidden Layer Neural Net

Hidden layer **pre-activation**:

$$\mathbf{a}(\mathbf{x}) = \mathbf{b}^{(1)} + \mathbf{W}^{(1)}\mathbf{x}$$

$$\left(\mathbf{a}(\mathbf{x})_i = b_i^{(1)} + \sum_j \mathbf{W}_{i,j}^{(1)} x_j \right)$$


Hidden layer **post-activation**:

$$\mathbf{h}(\mathbf{x}) = \sigma(\mathbf{a}(\mathbf{x}))$$

Output layer activation:

$$\mathbf{f}(\mathbf{x}) = o(b^{(2)} + \mathbf{w}^{(2)T} \mathbf{h}^{(1)}(\mathbf{x}))$$

Output activation function



Softmax output activation

In **multi-way classification**, we need multiple outputs (1 per class)

Natural: model calculates conditional probabilities $P(\text{output} = c | \mathbf{x})$

Softmax activation function at the output

$$\mathbf{o}(\mathbf{a}) = \text{softmax}(\mathbf{a}) = \left[\frac{\exp(a_1)}{\sum_c \exp(a_c)} \cdots \frac{\exp(a_C)}{\sum_c \exp(a_c)} \right]^\top$$

☯ strictly positive

☯ sums to one

Predict class with the highest estimated class conditional probability.

Multilayer Neural Net

Consider a network with L hidden layers.

Layer **pre-activations**:

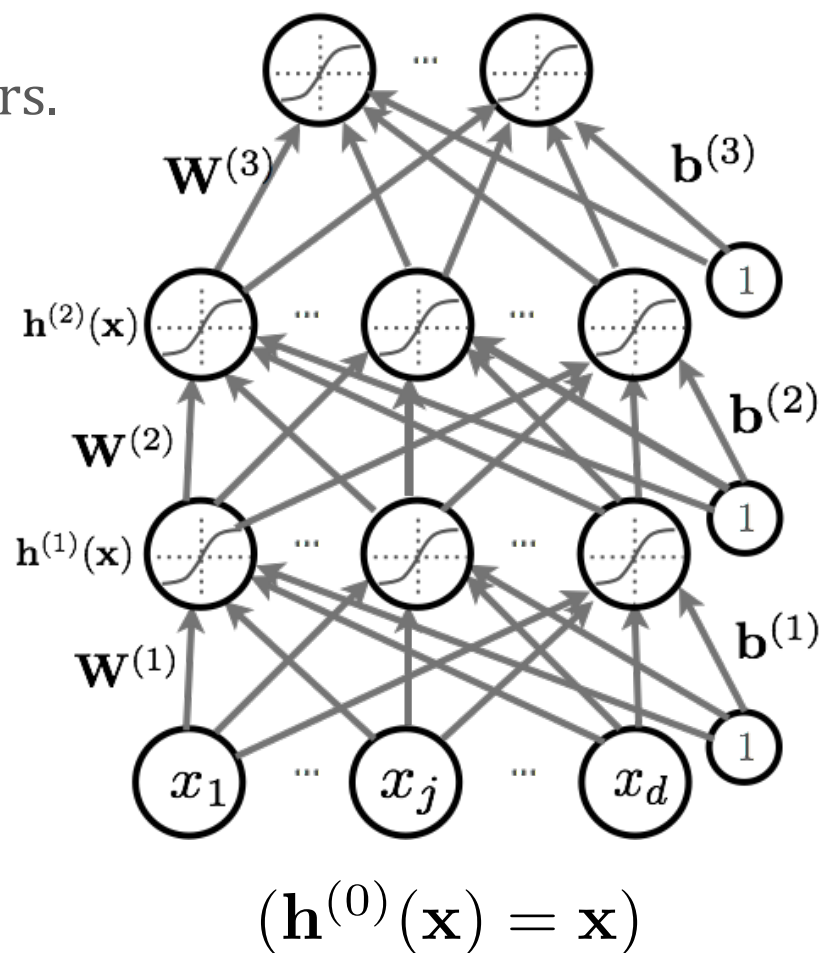
$$\mathbf{a}^{(k)}(\mathbf{x}) = \mathbf{b}^{(k)} + \mathbf{W}^{(k)} \mathbf{h}^{(k-1)}(\mathbf{x})$$

Hidden layer post-activations:

$$\mathbf{h}^{(k)}(\mathbf{x}) = \sigma(\mathbf{a}^{(k)}(\mathbf{x}))$$

Output layer activation:

$$\mathbf{h}^{(L+1)}(\mathbf{x}) = o(\mathbf{a}^{(L+1)}(\mathbf{x})) = \mathbf{f}(\mathbf{x})$$




Loss functions

Recall: typical approach is to minimize a **training** loss l over **predictors** \mathcal{F} :


$$\hat{f} = \operatorname{argmin}_{f \in \mathcal{F}} \sum_{(x,y): \text{training samples}} l(f(x), y)$$

Common losses:

l_2 : $l(f(x), y) = ||f(x) - y||^2$, more common for **regression**,
y can be vector or scalar



$$f(x) = \langle w, x \rangle$$



$$f(x) = \sum_i a_i x^i$$

Loss functions


Recall: typical approach is to minimize a **training** loss l over **predictors** \mathcal{F} :

$$\hat{f} = \operatorname{argmin}_{f \in \mathcal{F}} \sum_{(x,y): \text{training samples}} l(f(x), y)$$


Common losses:

l_2 : $l(f(x), y) = ||f(x) - y||^2$, more common for **regression**,
 y can be vector or scalar

0 – 1: $l(f(x), y) = 1_{f(x) \neq y}$, ideal loss for **classification**, but
poorly behaved for optimization



$$f(x) = \operatorname{sgn}(\langle w, x \rangle) \quad f(x) = \operatorname{sgn}\left(\sum_i a_i x^i\right)$$



Loss functions

Recall: typical approach is to minimize a **training** loss l over **predictors** \mathcal{F} :

$$\hat{f} = \operatorname{argmin}_{f \in \mathcal{F}} \sum_{(x,y): \text{training samples}} l(f(x), y)$$

Common losses:

l_2 : $l(f(x), y) = ||f(x) - y||^2$, more common for **regression**,
y can be vector or scalar

0 – 1: $l(f(x), y) = 1_{f(x) \neq y}$, ideal loss for **classification**, but
poorly behaved for optimization

Log-loss/ : cross entropy $l(f(x), y) = -\log f(x)_y$, for f using a **softmax** output layer,
well-behaved gradients

For softmax, $f(x)_c = P(\text{output} = c|x)$, so we **maximize** the log-probability of correct label. Generalizes naturally when y **not** deterministic fn of x in \mathcal{D} :

$$-\log f(x)_y = -\sum_c 1_{y=c} \log f(x)_c = -\sum_c 1_{y=c} \log P(\text{output} = c|x)$$

Taking expectation of y: $\mathbb{E}_{y|x} l(f(x), y) = -\mathbb{E}_{y|x} \log P(\text{output} = y | x)$

Basic optimization algorithm: stochastic gradient descent

Recall: typical approach is to minimize a **training** loss l over **predictors** \mathcal{F} :

$$\hat{f} = \operatorname{argmin}_{f \in \mathcal{F}} \sum_{(x,y): \text{training samples}} l(f(x), y)$$

Basic algorithm (Stochastic Gradient Descent)

Glossing over many details. Stay tuned.

- **Initialize:** $\theta_0 := \{W^{(1)}, b^{(1)}, \dots, W^{(L+1)}, b^{(L+1)}\}$
- For $t=1$ to T
 - Pick a uniformly random training example (x, y) :
 - Set $\theta_{t+1} = \theta_t - \eta \nabla_{\theta} l(f_{\theta}(x, y))$

Step size

*“Steepest” descent:
direction of most (local)
improvement*

*Neural nets:
gradients can be efficiently
calculated, using
backpropagation*

Supervised learning

Empirical risk minimization approach:
minimize a **training** loss l over a class of **predictors** \mathcal{F} :



$$\hat{f} = \operatorname{argmin}_{f \in \mathcal{F}} \sum_{(x,y): \text{training samples}} l(f(x), y)$$

Three pillars:

- (1) How expressive is the class \mathcal{F} ? (**Representational power**)
- (2) How do we minimize the training loss efficiently? (**Optimization**)
- (3) How does \hat{f} perform on unseen samples? (**Generalization**)

Expressivity


What do we mean by expressivity?




Expressive = functions in class can represent “complicated” functions

Linear classification

The arguably simplest class of classifiers is **linear**:




$$f(x) = \langle w, x \rangle$$



$$f(x) = \text{sgn}(\langle w, x \rangle)$$

How do we make classifiers “more expressive”?

One pervasive idea in machine learning (from kernels onward): train a linear classifier on a **feature embedding** of data.



$$f(x) = \sum_{i=0}^k a_i x^i$$

For instance, we can write

$f(x) = \langle a, \phi(x) \rangle$, where


$$a = (a_1, a_2, \dots, a_k)^T, \phi(x) = (1, x, x^2, \dots, x^k)$$

Hence, we first embed x via ϕ from \mathbb{R} into \mathbb{R}^k , and train a linear classifier on these new features.

How do we make classifiers “more expressive”?

By increasing degree we can increase expressiveness *a lot*:

For finite set of points $\{(x_1, y_1), \dots, (x_n, y_n)\}$, $x_i \in \mathbb{R}^d, y_i \in \mathbb{R}$ by **Lagrange's interpolation theorem**, we can find a polynomial p of degree $d - 1$, s.t. $\forall i, y_i = f(x_i)$




How do we make classifiers “more expressive”?

By increasing degree we can increase expressiveness *a lot*:

For any function f , we can **approximate it** on any compact set Ω by a sufficiently high degree polynomial: for every $\epsilon > 0$, $\exists p$ of sufficiently high degree, s.t.

$$\max_{x \in \Omega} |f(x) - p(x)| \leq \epsilon$$

(Stone-Weierstrass)



Vague intuition: think of Taylor series; near point x_0 , we have

$$f(x) = f(x_0) + f'(x_0)(x - x_0) + \frac{f''(x_0)}{2}(x - x_0)^2 + \dots$$

$$f(x) = \langle w, \phi(x) \rangle, \quad \phi(x) = (1, x, x^2, \dots), \quad w = (f(x_0), f'(x_0), \dots)$$

Lots of choices!

The name of the game in **kernel methods** was choose a good embedding ϕ we explored. Lots of latitude here:

Polynomial kernel (in d dim.):

$$\phi(x) = (1, x_1, x_2, \dots, x_d, x_1^2, x_1 x_2, \dots, x_d^2, \dots, x_d^k)$$

Gaussian kernel (in 1d.):

$$\phi(x) = e^{-\frac{x^2}{2\sigma^2}} \left(1, \sqrt{\frac{1}{1!}} \frac{x}{\sigma}, \sqrt{\frac{1}{2!}} \sqrt{\frac{1}{1!}} \left(\frac{x}{\sigma}\right)^2, \dots \right)$$

Choices of these kernels is closely related to something called the “kernel trick”, which allows for cheap computation of the **kernel** $\langle \phi(x), \phi(y) \rangle$. Beyond the scope of this course!

Lots of choices!

The name of the game in **kernel methods** was choose a good embedding ϕ we explored. Lots of latitude here:

Polynomial kernel (in d dim.):

$$\phi(x) = (1, x_1, x_2, \dots, x_d, x_1^2, x_1 x_2, \dots, x_d^2, \dots, x_d^k)$$

Gaussian kernel (in 1d.):

$$\phi(x) = e^{-\frac{x^2}{2\sigma^2}} \left(1, \sqrt{\frac{1}{1!}} \frac{x}{\sigma}, \sqrt{\frac{1}{2!}} \sqrt{\frac{1}{1!}} \left(\frac{x}{\sigma}\right)^2, \dots \right)$$

Choices of these kernels is closely related to something called the “kernel trick”, which allows for cheap computation of the **kernel** $\langle \phi(x), \phi(y) \rangle$. Beyond the scope of this course!


Part of the deep learning story




Object
detection



Image




vision features





Recognition


Audio
classification



Audio




audio features




Speaker
identification


Old school: hand-craft features




SIFT



Textons




HoG




RIFT


GIST




Old school: hand-craft features




Spectrogram




MFCC



Flux






ZCR




Rolloff


Old school: hand-craft features




Feature learning:
Can we automatically learn
useful features?



Flux

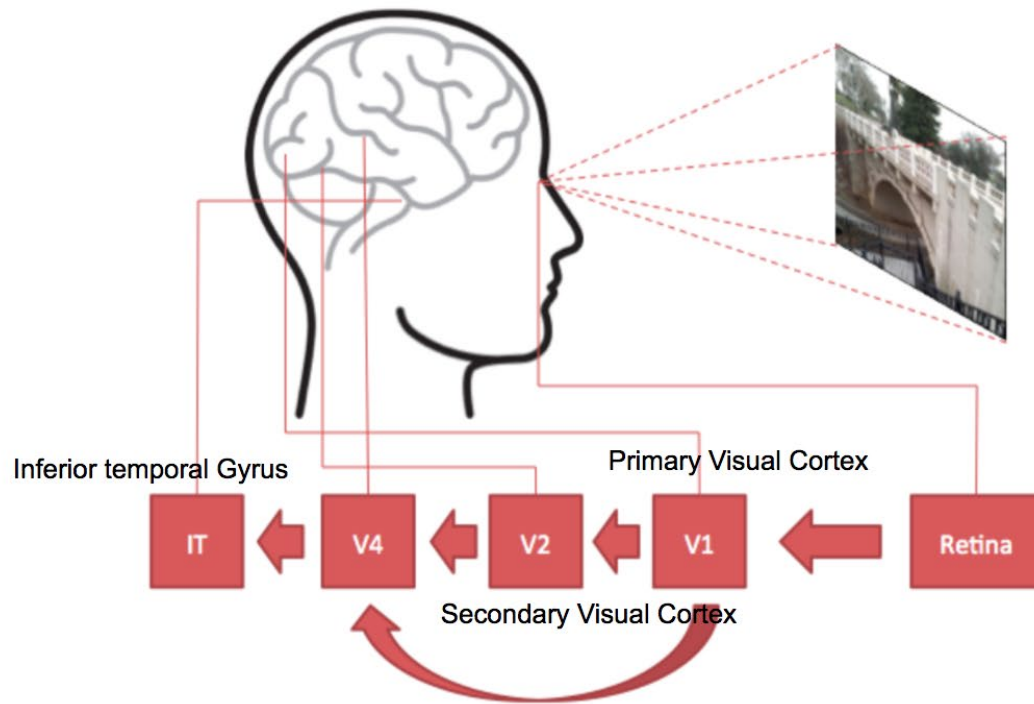


ZCR



Rolloff

Early inspirations from visual cortex



V1: Edge detection, etc.


V2: Extract simple visual properties (orientation, spatial frequency, color, etc)

V4: Detect object features of intermediate complexity

TI: Object recognition.

Image: Wang, Raj ["On the Origin of Deep Learning."](#)

What do deep networks learn?



Yosinski et al '15:

<http://yosinski.com/deepvis>

Q: What “patterns” do neurons respond to?

A: From random start, do gradient descent to find an input for which neuron activation* is high.



*: This produces completely unrecognizable images – they are regularized w/ an image prior.

What do deep networks learn?



Pirate Ship

Rocking Chair



Yosinski et al '15:

<http://yosinski.com/deepvis>

Q: What “patterns” do neurons respond to?


A: From random start, do gradient descent to find an input for which neuron activation* is high.

*: This produces completely unrecognizable images – they are regularized w/ an image prior.

“Universal” expressivity of neural networks

(1): Neural networks are **universal approximators**: given any Lipschitz $f: \mathbb{R}^d \rightarrow \mathbb{R}$, a **shallow** (3-layer) neural network with $\sim \left(\frac{1}{\epsilon}\right)^d$ neurons can approximate it to within ϵ error.

“curse of dimensionality”



Universal approximation I: Lipschitz function are approximable

Recall, a function $f: [0,1]^d \rightarrow \mathbb{R}$ is **L-Lipschitz** (in an l_∞ sense) if:
 $\forall x, y \in [0,1]^d, |f(x) - f(y)| \leq L \max_{i \in [d]} |x_i - y_i|$

First, we show neural networks are **universal approximators**: given any Lipschitz function $f: [0,1]^d \rightarrow \mathbb{R}$, a **shallow** (3-layer) neural network with $\sim \left(\frac{1}{\epsilon}\right)^d$ neurons can approximate it to within ϵ error.


Theorem: For any **L-Lipschitz** function $f: [0,1]^d \rightarrow \mathbb{R}$, there is a **3-layer** neural network \hat{f} with $O\left(d \left(\frac{L}{\epsilon}\right)^d\right)$ ReLU neurons, s.t.

$$\int_{[0,1]^d} |f(x) - \hat{f}(x)| dx \leq \epsilon$$

 l_1 error

Universal approximation I: Proof intuition

Part 1: using Lipschitzness, we can “query” the values of function f approximately by querying its values on a fine grid.



$$|f(y) - f(x)| \leq L\delta$$

Part 2: we can approximate f as linear combination of “queries”.

$$f(x) \approx \sum_{\text{cells } C_i} 1_{x \in C_i} f(x_i)$$

$$f(x) \approx \langle w, \phi(x) \rangle, \quad \phi(x) = (1(x \in C_i))_i \\ w = (f(x_i))_i$$


Part 3: Approximate the indicators using ReLUs

Universal approximation I:

Part 1, formally

Lemma: Let $f: [0,1]^d \rightarrow \mathbb{R}$ be L -Lipschitz and $P = (C_1, C_2, \dots, C_N)$ a partition of $[0,1]^d$ into cells of side lengths at most δ . Consider any set $(x_1, x_2, \dots, x_N), x_i \in C_i$. Then:

$$\sup_{i \in N} \sup_{y \in C_i} |f(y) - f(x_i)| \leq L\delta$$



Proof: By Lipschitzness, we have

$$\begin{aligned} \forall i, y \in C_i: |f(y) - f(x_i)| &\leq L \max_{i \in [d]} |y - x_i| \\ &\leq L\delta \quad \square \end{aligned}$$

Universal approximation I:

Part 2, formally

Lemma: Let $f: [0,1]^d \rightarrow \mathbb{R}$ be 1-Lipschitz, $P = (C_1, C_2, \dots, C_N)$ a partition of $[0,1]^d$ into rectangles of side lengths at most δ , and a set (x_1, x_2, \dots, x_N) , $x_i \in C_i$. Then,

$$g(x) = \sum_{i=1}^N 1_{x \in C_i} f(x_i) \text{ satisfies } \sup_{x \in [0,1]^d} |f(x) - g(x)| \leq L \delta$$

Proof: Let $x \in C_i$. Then, $1_{x \in C_i} = 1$, and $1_{x \in C_j} = 0$ for $j \neq i$.

So, $g(x) = f(x_i)$.

By Lemma 1, $|f(x) - g(x)| = |f(x) - f(x_i)| \leq L \delta \quad \square$

Universal approximation I: approximating indicators of cells

Lemma: Let $C \subseteq \mathbb{R}^d$ be a cell, namely $C = \{x: x \in [l_i, r_i], i \in d\}$. Then, there exists a 2-layer network $\tilde{h}(x)$ of size $O(d)$ and ReLU activation, s.t. $\int_{x \in [0,1]^d} |\tilde{h}(x) - 1(x \in [l_i, r_i], i \in d)| dx \rightarrow 0$

Proof: First, write indicator for cell as:

*For any $\gamma > 0$,
we'll take γ small*

$$1(x \in [l_i, r_i], i \in d) = 1 \left(\sum_{i=1}^d (1(x_i \geq l_i) + 1(x_i \leq r_i)) \geq 2d - 1 + \gamma \right)$$

Why? x is in cell iff all the indicators $1(x_i \geq l_i) + 1(x_i \leq r_i)$ are on.

All these indicators are on iff they sum to $2d$.

(If at least one is off, they sum to $2d-1$)

If we can approximate indicators, we're all good!

Universal approximation I: approximating indicators of cells


Claim: For $\tau \geq 0$, $x \in \mathbb{R}$:

$$|1(x \geq 0) - (\sigma(\tau x) - \sigma(\tau x - 1))| = \begin{cases} \leq 1, & \text{if } 0 \leq x \leq 1/\tau \\ 0, & \text{otherwise} \end{cases}$$

Proof: Consider several cases:

Case 1, $x \leq 0$: $1(x \geq 0) = 0$ and $\sigma(\tau x) - \sigma(\tau x - 1) = 0$, so

$$1(x \geq 0) - (\sigma(\tau x) - \sigma(\tau x - 1)) = 0$$



Case 2, $x \geq 1/\tau$: $1(x \geq 0) = 1$ and $\sigma(\tau x) - \sigma(\tau x - 1) = 1$, so

$$1(x \geq 0) - (\sigma(\tau x) - \sigma(\tau x - 1)) = 0$$

Case 3, $0 \leq x \leq 1/\tau$: $1(x \geq 0) = 1$ and $\sigma(\tau x) - \sigma(\tau x - 1) = \tau x$, so

$$|1(x \geq 0) - (\sigma(\tau x) - \sigma(\tau x - 1))| = \tau x \leq 1$$

Universal approximation I: approximating indicators of cells

$$h(x) := 1 \left(\sum_{i=1}^d (1(x_i \geq l_i) + 1(x_i \leq r_i)) \geq 2d - 1 + \gamma \right)$$

Replace all indicators by difference of ReLUs. *What is the error?*

For brevity, let $\tilde{1}(x \geq 0) = \sigma(\tau x) - \sigma(\tau x - 1)$, for some τ we will choose.

Let
$$\tilde{h}(x) := 1 \left(\sum_{i=1}^d (\tilde{1}(x_i \geq l_i) + \tilde{1}(x_i \leq r_i)) \geq 2d - 1 + \gamma \right)$$

$$\tilde{\tilde{h}}(x) := \tilde{1} \left(\sum_{i=1}^d (\tilde{1}(x_i \geq l_i) + \tilde{1}(x_i \leq r_i)) \geq 2d - 1 + \gamma \right)$$

(Change the approximations “iteratively”.)

Universal approximation I: approximating indicators of cells

$$\begin{aligned} h(x) &:= 1\left(\sum_{i=1}^d (1(x_i \geq l_i) + 1(x_i \leq r_i)) \geq 2d - 1 + \gamma\right) \\ \tilde{h}(x) &:= 1\left(\sum_{i=1}^d \left(\tilde{1}(x_i \geq l_i) + \tilde{1}(x_i \leq r_i)\right) \geq 2d - 1 + \gamma\right) \\ \tilde{\tilde{h}}(x) &:= \tilde{1}\left(\sum_{i=1}^d \left(\tilde{1}(x_i \geq l_i) + \tilde{1}(x_i \leq r_i)\right) \geq 2d - 1 + \gamma\right) \end{aligned}$$

We have:

$$\begin{aligned} \int_{x \in [0,1]} |\tilde{\tilde{h}}(x) - h(x)| dx &= \int_{x \in [0,1]^d} |\tilde{\tilde{h}}(x) - \tilde{h}(x) + \tilde{h}(x) - h(x)| dx \\ &\stackrel{\text{Triangle inequality}}{\leq} \int_{x \in [0,1]^d} \left(|\tilde{\tilde{h}}(x) - \tilde{h}(x)| + |\tilde{h}(x) - h(x)| \right) dx \end{aligned}$$

Let's handle two terms one by one.

Universal approximation I: approximating indicators of cells

$$\begin{aligned}\tilde{h}(x) &:= 1 \left(\sum_{i=1}^d \left(\tilde{1}(x_i \geq l_i) + \tilde{1}(x_i \leq r_i) \right) \geq 2d - 1 + \gamma \right) \\ \tilde{\tilde{h}}(x) &:= \tilde{1} \left(\sum_{i=1}^d \left(\tilde{1}(x_i \geq l_i) + \tilde{1}(x_i \leq r_i) \right) \geq 2d - 1 + \gamma \right)\end{aligned}$$

Claim: $|1(x \geq 0) - (\sigma(\tau x) - \sigma(\tau x - 1))| = \begin{cases} \leq 1, & \text{if } 0 \leq x \leq 1/\tau \\ 0, & \text{otherwise} \end{cases}$

First: $\tilde{\tilde{h}}(x) - \tilde{h}(x) \neq 0$ only if $\exists i : x_i \in \left(l_i, l_i + \frac{1}{\tau}\right)$ or $x_i \in \left(r_i, r_i - \frac{1}{\tau}\right)$

(If $\tilde{\tilde{h}}(x) - \tilde{h}(x) \neq 0$, $\sum_i \tilde{1}(x_i \geq l_i) + \tilde{1}(x_i \leq r_i) - (2d - 1) \in [\gamma, \gamma + \frac{1}{\tau})$, and if condition above isn't satisfied, $\tilde{1} = 1$, so $\sum_i \tilde{1}(x_i \geq l_i) + \tilde{1}(x_i \leq r_i) - (2d - 1)$ is integer, so cannot belong to interval for small enough $\gamma > 0$.)

Measure of such x 's is bdd by $\sum_i \left(\int_{x \in [0,1]^d} 1 \, dx + \int_{x \in [0,1]^d} 1 \, dx \right) \leq \frac{2d}{\tau}$, so

$$\int_{x \in [0,1]^d} |\tilde{\tilde{h}}(x) - \tilde{h}(x)| \leq \frac{2d}{\tau} \quad \begin{matrix} x_i \in \left(l_i, l_i + \frac{1}{\tau}\right) \\ x_i \in \left(r_i, r_i - \frac{1}{\tau}\right) \end{matrix}$$

Universal approximation I: approximating indicators of cells

$$h(x) := 1(\sum_{i=1}^d (1(x_i \geq l_i) + 1(x_i \leq r_i)) \geq 2d - 1 + \gamma)$$

$$\tilde{h}(x) := 1(\sum_{i=1}^d (\tilde{1}(x_i \geq l_i) + \tilde{1}(x_i \leq r_i)) \geq 2d - 1 + \gamma)$$

Claim: $|1(x \geq 0) - (\sigma(\tau x) - \sigma(\tau x - 1))| = \begin{cases} \leq 1, & \text{if } 0 \leq x \leq 1/\tau \\ 0, & \text{otherwise} \end{cases}$

Second: $\int_{x \in [0,1]} |\tilde{h}(x) - h(x)| dx$

Indicators are equal if inputs are equal

$$\leq \int_{x \in [0,1]^d} \left| 1 \left(\sum_{i=1}^d (1(x \geq l_i) + 1(x \leq r_i)) \neq \sum_{i=1}^d (\tilde{1}(x \geq l_i) + \tilde{1}(x \leq r_i)) \right) \right| dx$$

$$\leq \int_{x \in [0,1]^d} \sum_{i=1}^d 1(1(x \geq l_i) \neq \tilde{1}(x \geq l_i)) + \sum_{i=1}^d 1(1(x \leq r_i) \neq \tilde{1}(x \leq r_i)) dx$$

By Claim $\leq 2d/\tau$

Universal approximation I: approximating indicators of cells

$$\begin{aligned}h(x) &:= 1\left(\sum_{i=1}^d (1(x_i \geq l_i) + 1(x_i \leq r_i)) \geq 2d - 1 + \gamma\right) \\ \tilde{h}(x) &:= 1\left(\sum_{i=1}^d \left(\tilde{1}(x_i \geq l_i) + \tilde{1}(x_i \leq r_i)\right) \geq 2d - 1 + \gamma\right) \\ \tilde{\tilde{h}}(x) &:= \tilde{1}\left(\sum_{i=1}^d \left(\tilde{1}(x_i \geq l_i) + \tilde{1}(x_i \leq r_i)\right) \geq 2d - 1 + \gamma\right)\end{aligned}$$

Putting together, we have:

$$\begin{aligned}\int_{x \in [0,1]^d} |\tilde{\tilde{h}}(x) - h(x)| dx &= \int_{x \in [0,1]^d} |\tilde{\tilde{h}}(x) - \tilde{h}(x) + \tilde{h}(x) - h(x)| dx \\ &\leq 4d/\tau\end{aligned}$$

Also, $\tilde{\tilde{h}}(x)$ is a 2-layer net with ReLU activations and $O(d)$ nodes!

Universal approximation I: Putting everything together

By Part 1+2,
$$\sup_{x \in [0,1]^d} |f(x) - \sum_{i=1}^N 1_{x \in C_i} f(x_i)| \leq L \delta$$

Moreover, the number of cells N can be bounded by $\left(\frac{1}{\delta}\right)^d$

By indicator approximation: can approximate arbitrarily well by taking $\tau \rightarrow \infty$ with a 2-layer ReLU net.

Combining the above two points, we get a $\left(\frac{1}{\delta}\right)^d$ -sized 3-layer net s.t.

$$\int_{[0,1]^d} |f(x) - \hat{f}(x)| dx \leq L \delta$$

Taking $\delta = \frac{\epsilon}{L}$, the theorem follows.

Parting thoughts

All results we proved are **existential**: they prove that a good approximator exists. Finding one efficiently (much less so using gradient descent) is a different matter.

The choices of non-linearities are usually very **flexible**: most results of the type we saw can be re-proven using different non-linearities. (Examples in homework.)

Many other results of similar flavor. For instance, there are also results that deep, but narrow networks are universal approximators.