

10707

Deep Learning: Spring 2021

Andrej Risteski

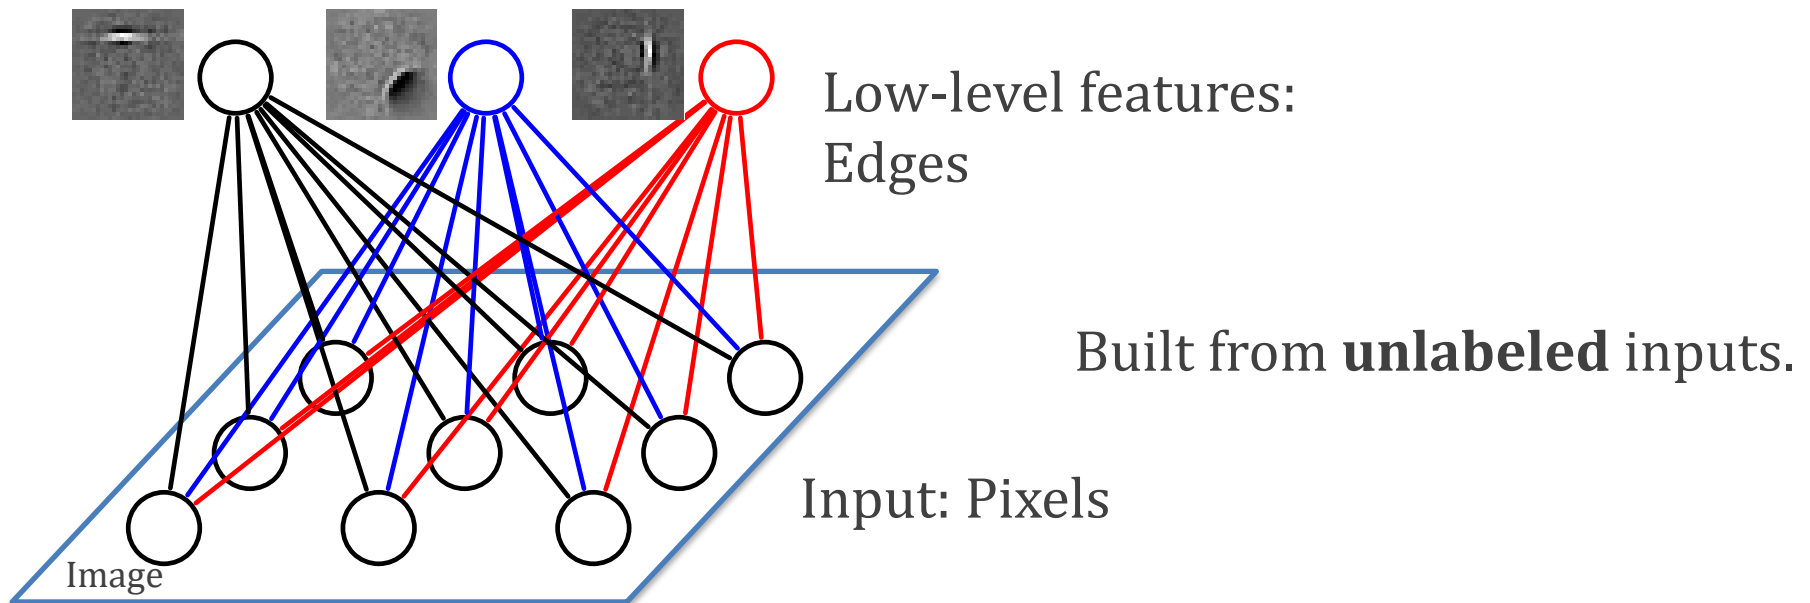
Machine Learning Department

Lecture 13:

More applications of
variational methods:
DBNs, VQ-VAEs, NVAEs

Part I: Learning Deep Belief Networks (DBNs)

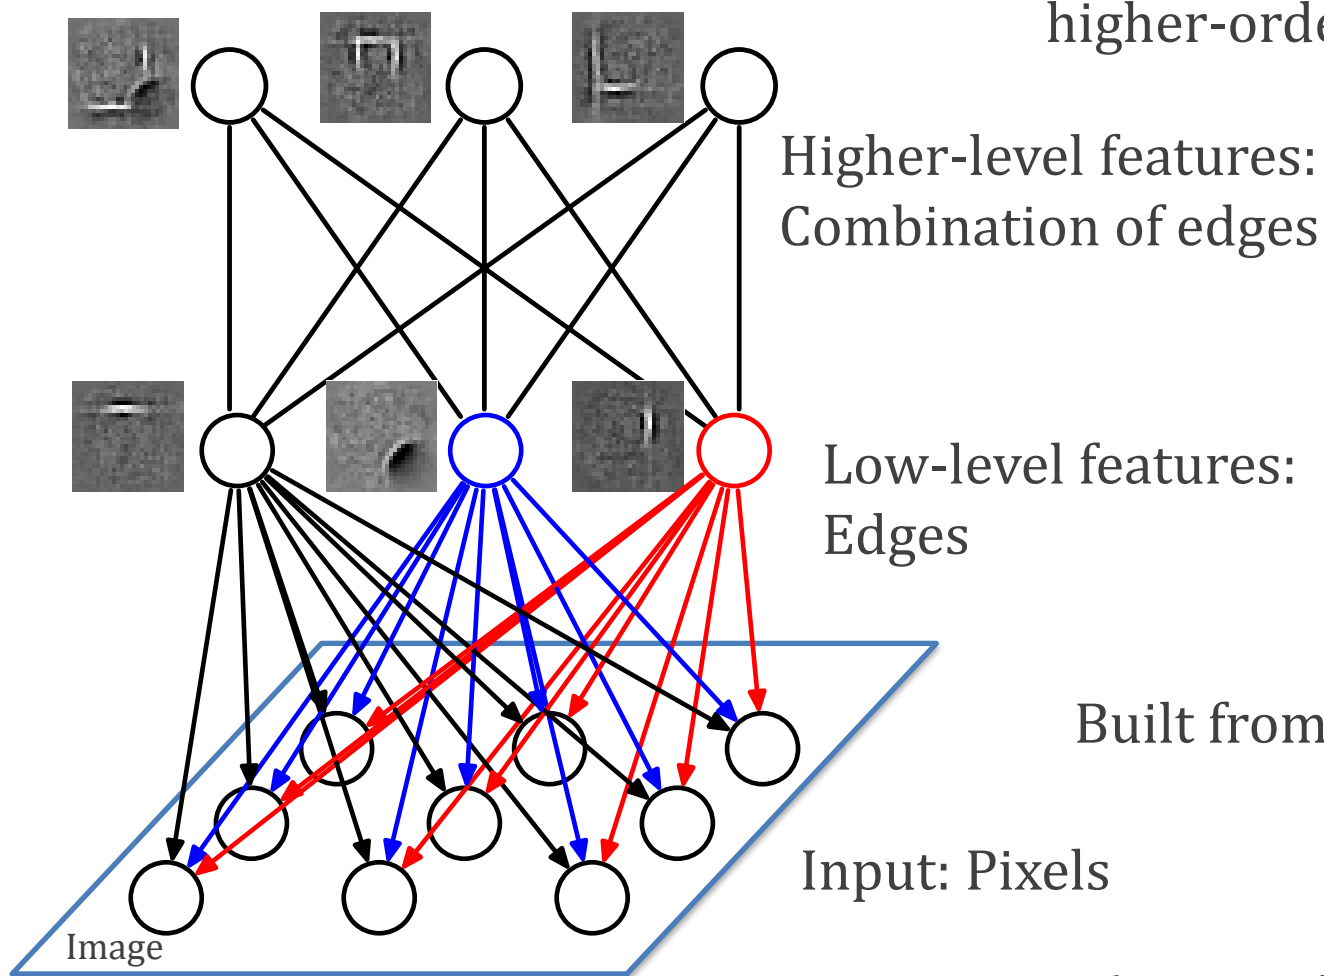
Deep Belief Network



(Hinton et.al. Neural Computation 2006)

Deep Belief Network

Internal representations capture higher-order statistical structure



Built from **unlabeled** inputs.


(Hinton et.al. Neural Computation 2006)

Deep Belief Network

- it is a generative model that mixes undirected and directed connections between variables
- top 2 layers' distribution $p(\mathbf{h}^{(2)}, \mathbf{h}^{(3)})$ is an RBM!
- other layers form a Bayesian network with conditional distributions:

$$p(h_j^{(1)} = 1 | \mathbf{h}^{(2)}) = \text{sigm}(\mathbf{b}^{(1)} + \mathbf{W}^{(2)\top} \mathbf{h}^{(2)})$$

$$p(x_i = 1 | \mathbf{h}^{(1)}) = \text{sigm}(\mathbf{b}^{(0)} + \mathbf{W}^{(1)\top} \mathbf{h}^{(1)})$$



Deep Belief Network

The **joint distribution** of a DBN is as follows

$$p(\mathbf{x}, \mathbf{h}^{(1)}, \mathbf{h}^{(2)}, \mathbf{h}^{(3)}) = p(\mathbf{h}^{(2)}, \mathbf{h}^{(3)}) p(\mathbf{h}^{(1)} | \mathbf{h}^{(2)}) p(\mathbf{x} | \mathbf{h}^{(1)})$$

where

$$p(\mathbf{h}^{(2)}, \mathbf{h}^{(3)}) = \exp \left(\mathbf{h}^{(2)\top} \mathbf{W}^{(3)} \mathbf{h}^{(3)} + \mathbf{b}^{(2)\top} \mathbf{h}^{(2)} + \mathbf{b}^{(3)\top} \mathbf{h}^{(3)} \right) / Z$$

$$p(\mathbf{h}^{(1)} | \mathbf{h}^{(2)}) = \prod_j p(h_j^{(1)} | \mathbf{h}^{(2)})$$

$$p(\mathbf{x} | \mathbf{h}^{(1)}) = \prod_i p(x_i | \mathbf{h}^{(1)})$$

(I realize this looks odd.)


DBN Layer-wise Training

Approximate
Inference

$$Q(\mathbf{h}^3|\mathbf{h}^2)$$

$$Q(\mathbf{h}^2|\mathbf{h}^1)$$

$$Q(\mathbf{h}^1|\mathbf{v})$$



Generative
Process

$$P(\mathbf{h}^2, \mathbf{h}^3)$$

$$P(\mathbf{h}^1|\mathbf{h}^2)$$


$$P(\mathbf{v}|\mathbf{h}^1)$$

$Q(h^t|h^{t-1}), P(h^{t-1}|h^t)$ are product distributions, s.t.:

$$Q\left((h^t)_j = 1 \mid h^{t-1}\right) = \frac{1}{1 + \exp(W_{t,j} h^{t-1})} \quad P\left((h^{t-1})_j = 1 \mid h^t\right) = \frac{1}{1 + \exp((h^t)^T W_{.,t})}$$


DBN Layer-wise Training

- Learn an RBM with an input layer $v=x$ and a hidden layer h .



DBN Layer-wise Training


- Learn an RBM with an input layer $\mathbf{v}=\mathbf{x}$ and a hidden layer \mathbf{h} .
- Treat inferred values $Q(\mathbf{h}^1|\mathbf{v}) = P(\mathbf{h}^1|\mathbf{v})$ as the data for training 2nd-layer RBM.
- Learn and freeze 2nd layer RBM.



DBN Layer-wise Training

- Learn an RBM with an input layer $\mathbf{v}=\mathbf{x}$ and a hidden layer \mathbf{h} .
- Treat inferred values $Q(\mathbf{h}^1|\mathbf{v}) = P(\mathbf{h}^1|\mathbf{v})$ as the data for training 2nd-layer RBM.
- Learn and freeze 2nd layer RBM.
- Proceed to the next layer.

Unsupervised Feature Learning.




Where does this training come from??

Variational intuitions

Let's write the marginal $p(\mathbf{x})$ in terms of the **Gibbs variational principle**.

Recall, we have:



For every distribution $q(\mathbf{h}^{(1)}|\mathbf{x})$:


$$\begin{aligned} \log p(\mathbf{x}) &\geq \sum_{\mathbf{h}^{(1)}} q(\mathbf{h}^{(1)}|\mathbf{x}) \log p(\mathbf{x}, \mathbf{h}^{(1)}) \\ &\quad - \sum_{\mathbf{h}^{(1)}} q(\mathbf{h}^{(1)}|\mathbf{x}) \log q(\mathbf{h}^{(1)}|\mathbf{x}) \end{aligned}$$

Equality is attained if $q(\mathbf{h}^{(1)}|\mathbf{x}) = p(\mathbf{h}^{(1)}|\mathbf{x})$.

Variational intuitions

Let's write the marginal $p(\mathbf{x})$ in terms of the **Gibbs variational principle**.

Recall, we have:




For every distribution $q(\mathbf{h}^{(1)}|\mathbf{x})$:

$$\log p(\mathbf{x}) \geq \sum_{\mathbf{h}^{(1)}} q(\mathbf{h}^{(1)}|\mathbf{x}) \log p(\mathbf{x}, \mathbf{h}^{(1)}) - \sum_{\mathbf{h}^{(1)}} q(\mathbf{h}^{(1)}|\mathbf{x}) \log q(\mathbf{h}^{(1)}|\mathbf{x})$$

*The idea will be to add layers, s.t. we improve the **variational bound (i.e. the right-hand side)***

Variational intuitions

adding 2nd layer means
untying the parameters



$$\begin{aligned}\log p(\mathbf{x}) &\geq \sum_{\mathbf{h}^{(1)}} q(\mathbf{h}^{(1)}|\mathbf{x}) \left(\log p(\mathbf{x}|\mathbf{h}^{(1)}) + \log p(\mathbf{h}^{(1)}) \right) \\ &\quad - \sum_{\mathbf{h}^{(1)}} q(\mathbf{h}^{(1)}|\mathbf{x}) \log q(\mathbf{h}^{(1)}|\mathbf{x})\end{aligned}$$


- When adding a second layer, we model $p(\mathbf{h}^{(1)})$ using a separate set of parameters

- they are the parameters of the RBM involving $\mathbf{h}^{(1)}$ and $\mathbf{h}^{(2)}$
- $p(\mathbf{h}^{(1)})$ is now the marginalization of the second hidden layer

$$p(\mathbf{h}^{(1)}) = \sum_{\mathbf{h}^{(2)}} p(\mathbf{h}^{(1)}, \mathbf{h}^{(2)})$$

Variational intuitions

adding 2nd layer means
untying the parameters



$$\log p(\mathbf{x}) \geq \sum_{\mathbf{h}^{(1)}} q(\mathbf{h}^{(1)}|\mathbf{x}) \left(\log p(\mathbf{x}|\mathbf{h}^{(1)}) + \log p(\mathbf{h}^{(1)}) \right) - \sum_{\mathbf{h}^{(1)}} q(\mathbf{h}^{(1)}|\mathbf{x}) \log q(\mathbf{h}^{(1)}|\mathbf{x})$$

- we can train the parameters of the **bound**. This is equivalent to maximizing the terms that are constant:

$$- \sum_{\mathbf{h}^{(1)}} q(\mathbf{h}^{(1)}|\mathbf{x}) \log p(\mathbf{h}^{(1)})$$

Layerwise training
improves variational
lower bound

- this is like training an RBM on data **generated** from $q(\mathbf{h}^{(1)}|\mathbf{x})$!

Stacking the layers




This is where the RBM stacking procedure comes from:

- **idea:** improve prior on last layer by adding another hidden layer

$$p(\mathbf{h}^{(1)}, \mathbf{h}^{(2)}) = p(\mathbf{h}^{(1)} | \mathbf{h}^{(2)}) \sum_{\mathbf{h}^{(3)}} p(\mathbf{h}^{(2)}, \mathbf{h}^{(3)})$$

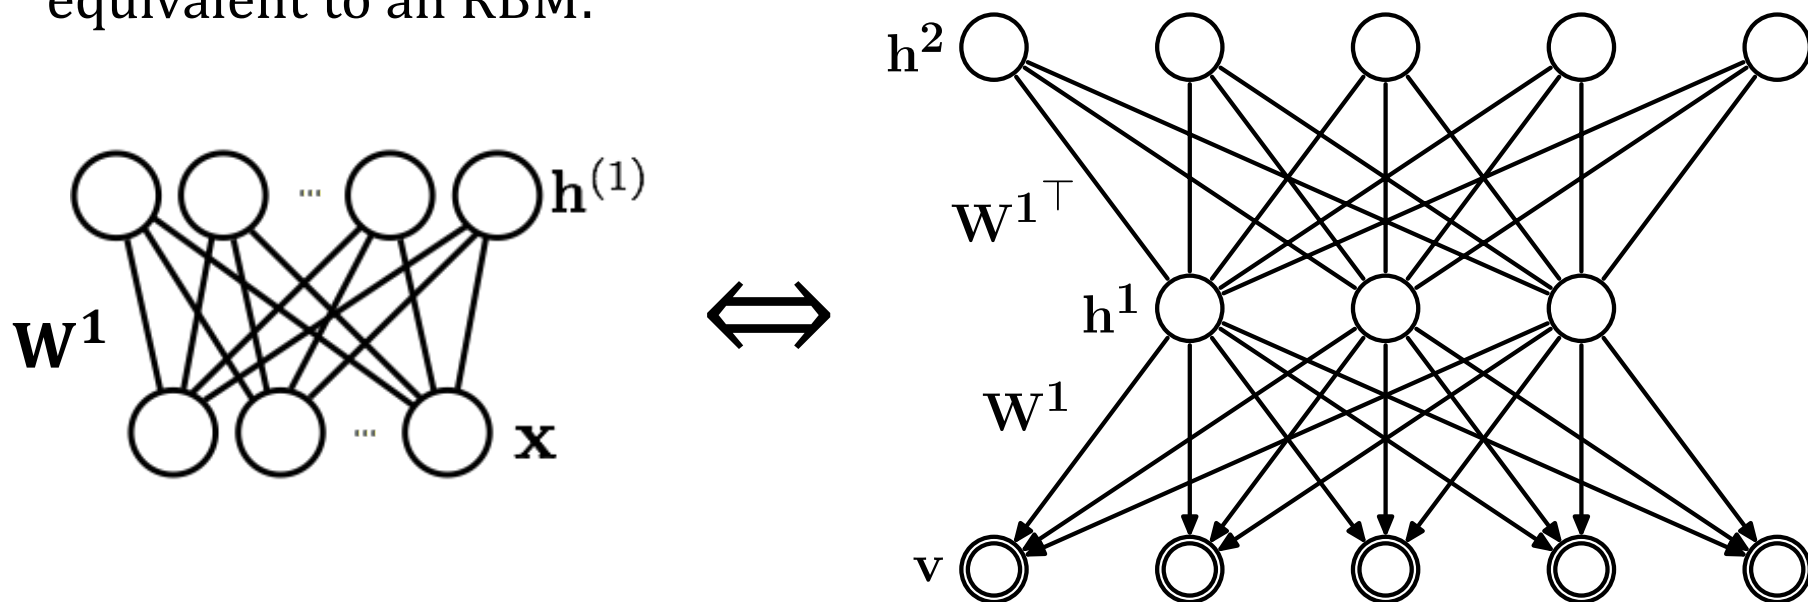
$$p(\mathbf{x}, \mathbf{h}^{(1)}) = p(\mathbf{x} | \mathbf{h}^{(1)}) \sum_{\mathbf{h}^{(2)}} p(\mathbf{h}^{(1)}, \mathbf{h}^{(2)})$$

$$p(\mathbf{x}) = \sum_{\mathbf{h}^{(1)}} p(\mathbf{x}, \mathbf{h}^{(1)})$$



Improvement at initialization: weight-tied DBN is equivalent to a RBM

Observation: a two-layer DBN with appropriately tied weights is equivalent to an RBM:



Formal proof is a little annoying. Intuition:

- Gibbs sampling converges to model distribution in first case.
- Gibbs sampling on top two layers, plus one last sample of x given $h^{(1)}$ converges to model distribution in second.
- The steps in these two random walks are **exactly** the same.

Improvement at initialization: weight-tied DBN is equivalent to a RBM

adding 2nd layer means
untying the parameters

$$\begin{aligned} \log p(\mathbf{x}) \geq & \sum_{\mathbf{h}^{(1)}} q(\mathbf{h}^{(1)}|\mathbf{x}) \left(\log p(\mathbf{x}|\mathbf{h}^{(1)}) + \log p(\mathbf{h}^{(1)}) \right) \\ & - \sum_{\mathbf{h}^{(1)}} q(\mathbf{h}^{(1)}|\mathbf{x}) \log q(\mathbf{h}^{(1)}|\mathbf{x}) \end{aligned}$$

- for $q(\mathbf{h}^{(1)}|\mathbf{x})$ we use **the posterior of the first layer RBM**.
- by initializing the weights of the second layer RBM as the transpose of the first layer weights, **the bound is initially tight!** (As we showed, a 2-layer DBN with tied weights is equivalent to a 1-layer RBM)
- Need not keep being tight:
as $p(\mathbf{h}^{(1)})$ changes, so does $p(\mathbf{h}^{(1)}|\mathbf{x})$, and so does the KL to $q(\mathbf{h}^{(1)}|\mathbf{x})$

Deep Belief Networks

This process of adding layers can be repeated recursively

- we obtain the greedy layer-wise pre-training procedure for neural networks

We now see that this procedure corresponds to maximizing a bound on the likelihood of the data in a DBN

- in theory, if our approximation $q(\mathbf{h}^{(1)}|\mathbf{x})$ is very far from the true posterior, the bound might be very loose
- this only means we might not be improving the true likelihood
- we might still be extracting better features!

Fine-tuning is done by the Up-Down algorithm



- A fast learning algorithm for deep belief nets. Hinton, Teh, Osindero, 2006.

Sampling from DBNs

- To sample from the DBN model:


$$P(\mathbf{v}, \mathbf{h}^1, \mathbf{h}^2, \mathbf{h}^3) = P(\mathbf{v}|\mathbf{h}^1)P(\mathbf{h}^1|\mathbf{h}^2)P(\mathbf{h}^2, \mathbf{h}^3)$$

- Sample \mathbf{h}^2 using alternating Gibbs sampling from RBM.
- Sample lower layers using sigmoid belief network.




Learned Features

1st-layer features




2nd-layer features

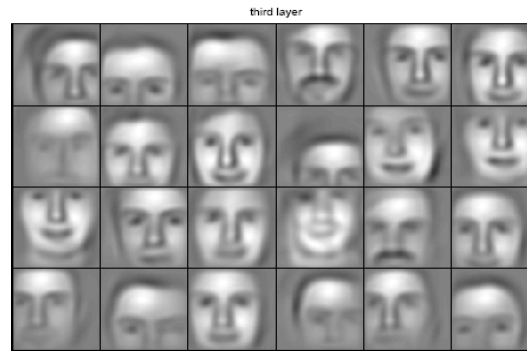


Learning Part-based Representation

Convolutional DBN



Faces




Groups of parts.



second layer





Object Parts




Trained on face images.

Learning Part-based Representation



Faces






Cars




Elephants




Chairs




Learning Part-based Representation



Groups of parts.



Class-specific object parts



Trained from multiple classes (cars, faces, motorbikes, airplanes).

Part II: Variants of VAEs

(VQ-VAEs, NVAEs)

VQ-VAE, Oord et al '17, Razavi et al '19



Figure from Razavi et al '19

Basic idea: discrete latent space

Idea: perform k-means on the recovered latent vectors to discretize

*Map the latent vectors to the
closest mean
(nearest-neighbor lookup):*

$$q(z = k|x) = \begin{cases} 1 & \text{for } k = \operatorname{argmin}_j \|z_e(x) - e_j\|_2, \\ 0 & \text{otherwise} \end{cases}$$

Latent embedding space

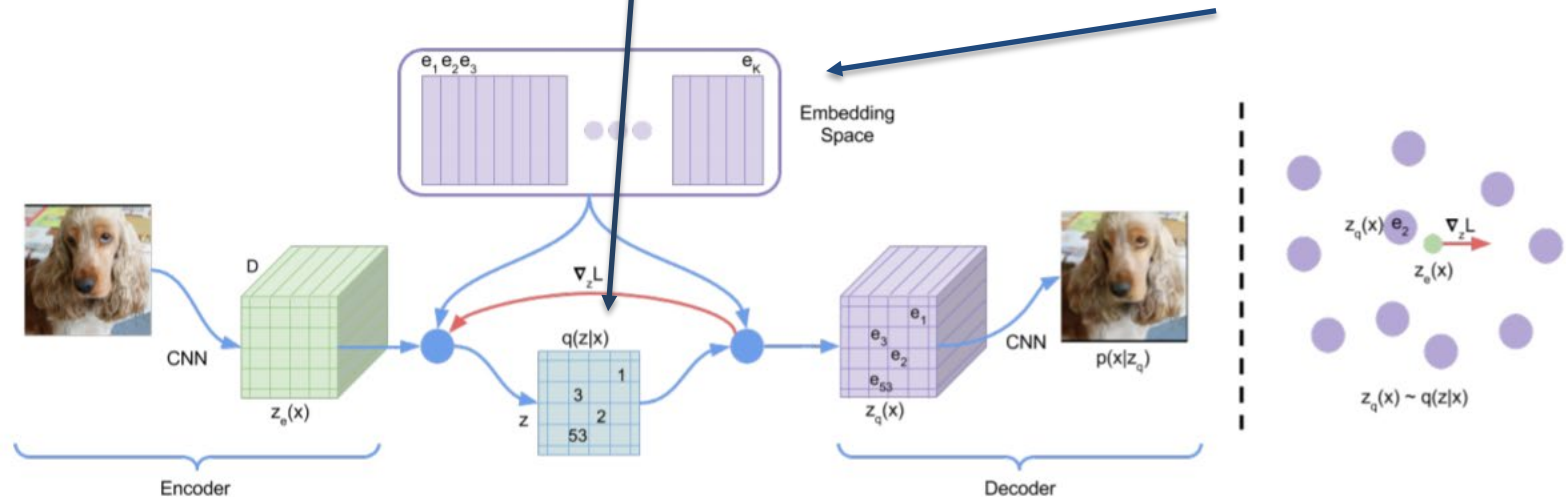


Figure 1: Left: A figure describing the VQ-VAE. Right: Visualisation of the embedding space. The output of the encoder $z(x)$ is mapped to the nearest point e_2 . The gradient $\nabla_z L$ (in red) will push the encoder to change its output, which could alter the configuration in the next forward pass.

The loss

Idea: perform k-means on the recovered latent vectors to discretize

*Map the latent vectors to the
closest mean*

(nearest-neighbor lookup):

$$q(z = k|x) = \begin{cases} 1 & \text{for } k = \operatorname{argmin}_j \|z_e(x) - e_j\|_2, \\ 0 & \text{otherwise} \end{cases}$$

Variational posterior $q(z|x)$ is a distribution over a domain of size K , and a **point mass**. Let's denote $z_q(x) = \operatorname{argmin}_j \|z_e(x) - e_j\|_2$.

Loss (first try):
$$L(\theta, e) = \mathbb{E}_x \left[\underbrace{\log(p_\theta(x|z_q))}_{\text{Reconstruction loss}} + \underbrace{KL(q(h|x) \| p(h))}_{\text{Regularization towards prior}} \right]$$

The authors drop the regularization term.

Problem: the mapping $z_e \rightarrow z_q$ involves argmin and is not differentiable.

The loss: straight-through estimator

Idea: perform k-means on the recovered latent vectors to discretize

Map the latent vectors to the closest mean

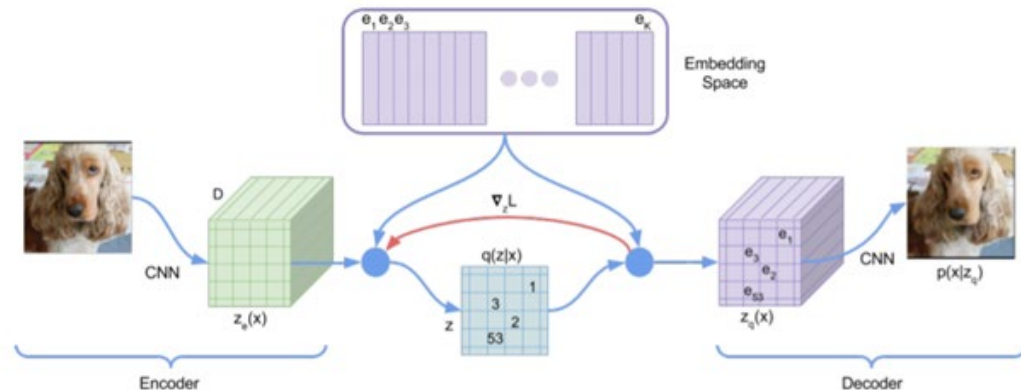
(nearest-neighbor lookup):

$$q(z = k|x) = \begin{cases} 1 & \text{for } k = \operatorname{argmin}_j \|z_e(x) - e_j\|_2, \\ 0 & \text{otherwise} \end{cases}$$

Loss (first try): $L(\theta, e) = \mathbb{E}_x[\log(p_\theta(x|z_q))]$

Problem: the mapping $z_e \rightarrow z_q$ involves argmin and is not differentiable.

Solution: use **straight-through estimator**; copy gradients from decoder input $z_q(x)$ to encoder output $z_e(x)$



The loss: quantization

Idea: perform k-means on the recovered latent vectors to discretize

*Map the latent vectors to the
closest mean*

(nearest-neighbor lookup):

$$q(z = k|x) = \begin{cases} 1 & \text{for } k = \operatorname{argmin}_j \|z_e(x) - e_j\|_2, \\ 0 & \text{otherwise} \end{cases}$$

Problem: the cluster means $\{e_j\}$ are not getting updated.

Loss (second try): $L(\theta, e) = \mathbb{E}_x \left[\underbrace{\log(p_\theta(x|z_q))}_{\text{Reconstruction loss}} + \underbrace{\left\| SG(z_e(x)) - z_q(x) \right\|^2}_{\text{Quantization loss}} \right]$

$SG(z_e(x))$: **stop-gradient operator**; identity at forward computation; has zero derivative, so argument doesn't get update at backward computation

This term only updates cluster means $\{e_j\}$!

The loss: quantization

Idea: perform k-means on the recovered latent vectors to discretize

*Map the latent vectors to the
closest mean*

(nearest-neighbor lookup):

$$q(z = k|x) = \begin{cases} 1 & \text{for } k = \operatorname{argmin}_j \|z_e(x) - e_j\|_2, \\ 0 & \text{otherwise} \end{cases}$$

Problem: the cluster means $\{e_j\}$ are not getting updated.

Loss (second try): $L(\theta, e) = \mathbb{E}_x \left[\underbrace{\log(p_\theta(x|z_q))}_{\text{Reconstruction loss}} + \underbrace{\left\| SG(z_e(x)) - z_q(x) \right\|^2}_{\text{Quantization loss}} \right]$

(Alternatively, this term can be rewritten as follows. If $z_{1,i}, \dots, z_{n_i,i}$ are the decoded samples closest to e_i , this term is $\sum_i \sum_{j=1}^{n_i} \|z_{j,i} - e_i\|^2$.

New e_i can just be set to $e_i := \sum_{j=1}^{n_i} z_{j,i}$.)

The loss: commitment penalty

Idea: perform k-means on the recovered latent vectors to discretize

*Map the latent vectors to the
closest mean*

(nearest-neighbor lookup):

$$q(z = k|x) = \begin{cases} 1 & \text{for } k = \operatorname{argmin}_j \|z_e(x) - e_j\|_2, \\ 0 & \text{otherwise} \end{cases}$$

Loss (third try):

$$L(\theta, e) = \mathbb{E}_x \left[\underbrace{\log(p_\theta(x|z_q))}_{\text{Reconstruction loss}} + \underbrace{\left\| SG(z_e(x)) - z_q(x) \right\|^2}_{\text{Quantization loss}} + \underbrace{\beta \left\| z_e(x) - SG(z_q(x)) \right\|^2}_{\text{Commitment loss}} \right]$$

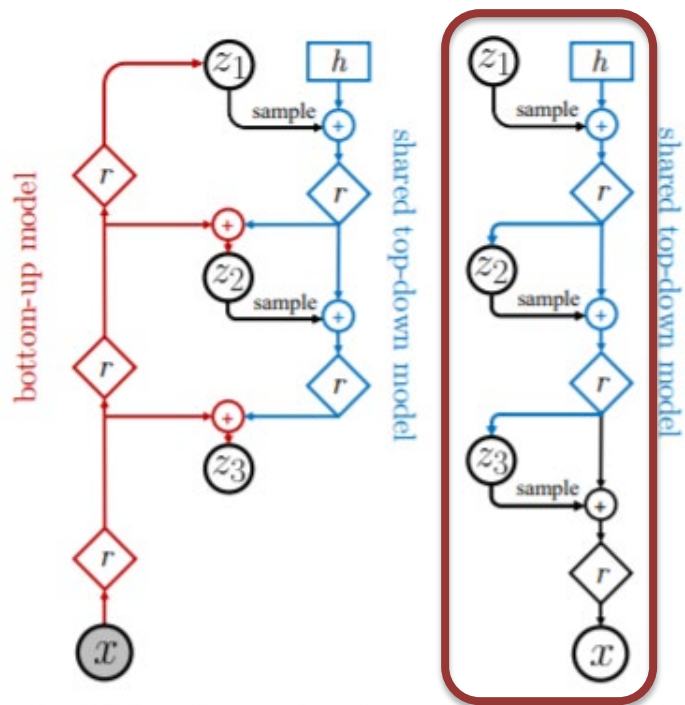
Authors add a **commitment loss**: this “regularizer” attempts to keep unquantized $z_e(x)$ close to current means $\{e_i\}$.

NVAE, Vahdat-Kautz '21



Figure from Vahdat-Kautz '21

Basic idea: careful changes in architecture



(a) Bidirectional Encoder (b) Generative Model

Figure 2: The neural networks implementing an encoder $q(\mathbf{z}|\mathbf{x})$ and generative model $p(\mathbf{x}, \mathbf{z})$ for a 3-group hierarchical VAE. \diamond_r denotes residual neural networks, \oplus denotes feature combination (e.g., concatenation), and \boxed{h} is a trainable parameter.

Main idea: hierarchical model for the generative and inference direction, with careful choice of architecture;

Generative model:

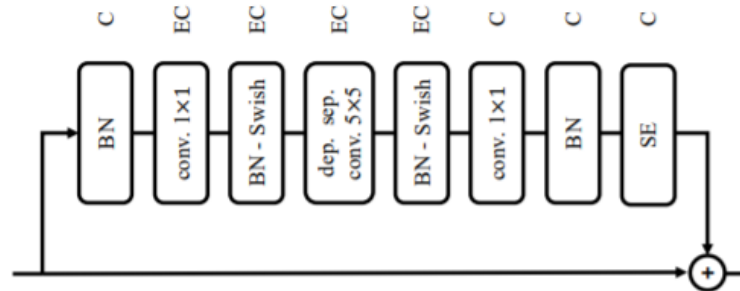
Use residual networks from layer to layer; the dimensions of the z 's gradually increase to gradually add more detail to image.

Basic idea: careful changes in architecture

Generative model:

Use residual networks from layer to layer; the dimensions of the z 's gradually increase to gradually add more detail to image.

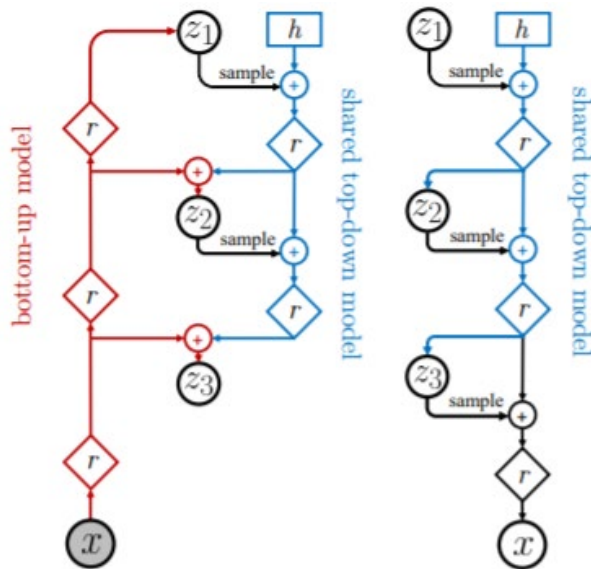
Hyperparameter	MNIST 28×28	CIFAR-10 32×32	ImageNet 32×32	CelebA 64×64	CelebA HQ 256×256	FFHQ 256×256
# latent variable scales	2	1	1	3	5	5
# groups in each scale	5, 10	30	28	5, 10, 20	4, 4, 4, 8, 16	4, 4, 4, 8, 16
spatial dims of z in each scale	4 ² , 8 ²	16 ²	16 ²	8 ² , 16 ² , 32 ²	8 ² , 16 ² , 32 ² , 64 ² , 128 ²	8 ² , 16 ² , 32 ² , 64 ² , 128 ²



(a) Residual Cell for NVAE Generative Model

Basic idea: careful changes in architecture

Encoder: weight tied w/ decoder for better behavior of KL term



(a) Bidirectional Encoder (b) Generative Model

Figure 2: The neural networks implementing an encoder $q(\mathbf{z}|\mathbf{x})$ and generative model $p(\mathbf{x}, \mathbf{z})$ for a 3-group hierarchical VAE. \diamond denotes residual neural networks, \oplus denotes feature combination (e.g., concatenation), and \boxed{h} is a trainable parameter.

Recall, there is a term $KL(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}))$ which will be large if q and p are far.

Weight tying: if we parametrize $p(\mathbf{z})$ s.t.

$$p(z_l^i | \mathbf{z}_{<l}) := \mathcal{N}(\mu_i(\mathbf{z}_{<l}), \sigma_i(\mathbf{z}_{<l}))$$

we parametrize $q(\mathbf{z}|\mathbf{x})$ correspondingly as

$$q(z_l^i | \mathbf{z}_{<l}, \mathbf{x}) := \mathcal{N}(\mu_i(\mathbf{z}_{<l}) + \Delta\mu_i(\mathbf{z}_{<l}, \mathbf{x}), \sigma_i(\mathbf{z}_{<l}) \cdot \Delta\sigma_i(\mathbf{z}_{<l}, \mathbf{x}))$$

(“relative to p ” parametrization)

Then, we have:

$$KL(q(z^i|\mathbf{x})||p(z^i)) = \frac{1}{2} \left(\frac{\Delta\mu_i^2}{\sigma_i^2} + \Delta\sigma_i^2 - \log \Delta\sigma_i^2 - 1 \right)$$